

Top 50 Large Language Model (LLM) Interview Questions

Hao Hoang - Follow me on LinkedIn for AI insights!

May 2025

Explore the key concepts, techniques, and challenges of Large Language Models (LLMs) with this comprehensive guide, crafted for AI enthusiasts and professionals preparing for interviews.

Introduction

Large Language Models (LLMs) are revolutionizing artificial intelligence, enabling applications from chatbots to automated content creation. This document compiles 50 essential interview questions, carefully curated to deepen your understanding of LLMs. Each question is paired with a detailed answer, blending technical insights with practical examples. Share this knowledge with your network to spark meaningful discussions in the AI community!

1 Question 1: What does tokenization entail, and why is it critical for LLMs?

Tokenization involves breaking down text into smaller units, or tokens, such as words, subwords, or characters. For example, "artificial" might be split into "art," "ific," and "ial." This process is vital because LLMs process numerical representations of tokens, not raw text. Tokenization enables models to handle diverse languages, manage rare or unknown words, and optimize vocabulary size, enhancing computational efficiency and model performance.

2 Question 2: How does the attention mechanism function in transformer models?

The attention mechanism allows LLMs to weigh the importance of different tokens in a sequence when generating or interpreting text. It computes similarity scores between query, key, and value vectors, using operations like dot products, to focus on relevant tokens. For instance, in "The cat chased the mouse," attention helps the model link "mouse" to "chased." This mechanism improves context understanding, making transformers highly effective for NLP tasks.

3 Question 3: What is the context window in LLMs, and why does it matter?

The context window refers to the number of tokens an LLM can process at once, defining its "memory" for understanding or generating text. A larger window, like 32,000 tokens, allows the model to consider more context, improving coherence in tasks like summarization. However, it increases computational costs. Balancing window size with efficiency is crucial for practical LLM deployment.

4 Question 4: What distinguishes LoRA from QLoRA in fine-tuning LLMs?

LoRA (Low-Rank Adaptation) is a fine-tuning method that adds low-rank matrices to a model's layers, enabling efficient adaptation with minimal memory overhead. QLoRA extends this by applying quantization (e.g., 4-bit precision) to further reduce memory usage while maintaining accuracy. For example, QLoRA can fine-tune a 70B-parameter model on a single GPU, making it ideal for resource-constrained environments.

5 Question 5: How does beam search improve text generation compared to greedy decoding?

Beam search explores multiple word sequences during text generation, keeping the top k candidates (beams) at each step, unlike greedy decoding, which selects only the most probable word. This approach, with $k = 5$, for instance, ensures more coherent outputs by balancing probability and diversity, especially in tasks like machine translation or dialogue generation.

6 Question 6: What role does temperature play in controlling LLM output?

Temperature is a hyperparameter that adjusts the randomness of token selection in text generation. A low temperature (e.g., 0.3) favors high-probability tokens, producing predictable outputs. A high temperature (e.g., 1.5) increases diversity by flattening the probability distribution. Setting temperature to 0.8 often balances creativity and coherence for tasks like storytelling.

7 Question 7: What is masked language modeling, and how does it aid pretraining?

Masked language modeling (MLM) involves hiding random tokens in a sequence and training the model to predict them based on context. Used in models like BERT, MLM fosters bidirectional understanding of language, enabling the model to grasp semantic